

Introduction:

This report represents the third assignment completed during my Data Science Internship. The aim of this task was to apply advanced data cleaning and preprocessing techniques to a large Airline Passenger Satisfaction dataset. The objective was to identify data inconsistencies and transform the dataset into an analysis-ready format using Python.

LIBRARIES IMPORT

```
[1] ✓
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

DATASET LOAD KARO

| df = pd.read_csv("./content/train.csv") df.head() | | | | | | | | | | | | | | | |
|--|----|--------|---------------|-------------------|----------------|-----------------|-----------------|-----------------------|-----------------------------------|-----|------------------------|------------------|------------------|------------------|----------------|
| Unnamed: 0 | id | Gender | Customer Type | Age | Type of Travel | Class | Flight Distance | Inflight wifi service | Departure/Arrival time convenient | ... | Inflight entertainment | On-board service | Leg room service | Baggage handling | Check-in serv. |
| 0 | 0 | 70172 | Male | Loyal Customer | 13 | Personal Travel | Eco Plus | 460 | 3 | 4 | ... | 5 | 4 | 3 | 4 |
| 1 | 1 | 5047 | Male | disloyal Customer | 25 | Business travel | Business | 235 | 3 | 2 | ... | 1 | 1 | 5 | 3 |
| 2 | 2 | 110028 | Female | Loyal Customer | 26 | Business travel | Business | 1142 | 2 | 2 | ... | 5 | 4 | 3 | 4 |
| 3 | 3 | 24026 | Female | Loyal Customer | 25 | Business travel | Business | 562 | 2 | 5 | ... | 2 | 2 | 5 | 3 |

Dataset Description:

The Airline Passenger Satisfaction dataset contains survey responses from airline customers. It includes features such as gender, age, travel class, flight distance, and delay information. Due to its size and complexity, this dataset was an excellent choice for practicing professional data cleaning techniques

INITIAL DIAGNOSTICS

```
df.shape
df.info()
df.describe(include='all')

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 103904 entries, 0 to 103903
Data columns (total 25 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        103904 non-null   int64  
 1   id               103904 non-null   int64  
 2   Gender            103904 non-null   object  
 3   Customer Type    103904 non-null   object  
 4   Age               103904 non-null   int64  
 5   Type of Travel   103904 non-null   object  
 6   Class              103904 non-null   object  
 7   Flight Distance   103904 non-null   int64  
 8   Inflight wifi service 103904 non-null   int64  
 9   Departure/Arrival time convenient 103904 non-null   int64  
 . . .
22  Departure Delay in Minutes      103904 non-null   int64  
23  Arrival Delay in Minutes       103904 non-null   float64 
24  satisfaction                  103904 non-null   object  
dtypes: float64(1), int64(19), object(5)
memory usage: 19.8+ MB
```

| | Unnamed: 0 | id | Gender | Customer Type | Age | Type of Travel | Class | Flight Distance | Inflight wifi service | Departure/Arrival time convenient | ... | Infl | entertainm |
|--------|---------------|---------------|--------|----------------|---------------|-----------------|----------|-----------------|-----------------------|-----------------------------------|-----|---------------|---------------|
| count | 103904.000000 | 103904.000000 | 103904 | 103904 | 103904.000000 | 103904 | 103904 | 103904.000000 | 103904.000000 | 103904.000000 | ... | 103904.000000 | 103904.000000 |
| unique | Nan | Nan | 2 | 2 | Nan | 2 | 3 | Nan | Nan | Nan | ... | Nan | 1 |
| top | Nan | Nan | Female | Loyal Customer | Nan | Business travel | Business | Nan | Nan | Nan | ... | Nan | 1 |
| freq | Nan | Nan | 52727 | 84923 | Nan | 71655 | 49665 | Nan | Nan | Nan | ... | Nan | 1 |
| mean | 51951.500000 | 64924.210502 | Nan | Nan | 39.379706 | Nan | Nan | 1189.448375 | 2.729683 | 3.060296 | ... | 3.358 | |
| std | 29994.645522 | 37463.812252 | Nan | Nan | 15.114964 | Nan | Nan | 997.147281 | 1.327829 | 1.525075 | ... | 1.332 | |
| min | 0.000000 | 1.000000 | Nan | Nan | 7.000000 | Nan | Nan | 31.000000 | 0.000000 | 0.000000 | ... | 0.000 | |
| 25% | 25975.750000 | 32533.750000 | Nan | Nan | 27.000000 | Nan | Nan | 414.000000 | 2.000000 | 2.000000 | ... | 2.000 | |

Missing Value Analysis:

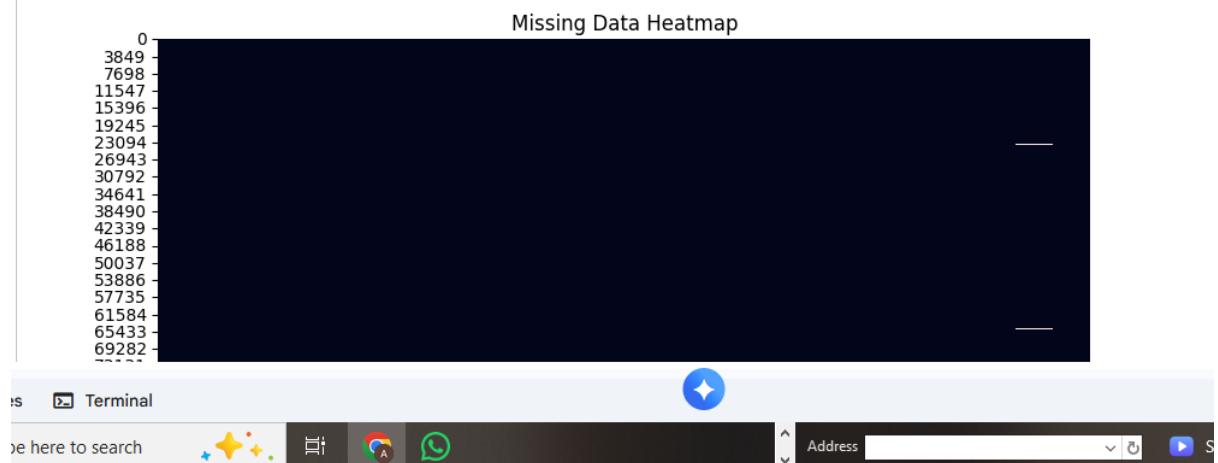
Missing values were identified in several columns including delay-related features. To better understand their distribution, tabular summaries and heatmap visualization were used.

MISSING VALUES ANALYSIS

| df.isnull().sum() | |
|-----------------------|---|
| | 0 |
| Unnamed: 0 | 0 |
| id | 0 |
| Gender | 0 |
| Customer Type | 0 |
| Age | 0 |
| Type of Travel | 0 |
| Class | 0 |
| Flight Distance | 0 |
| Inflight wifi service | 0 |

HeatMap

```
plt.figure(figsize=(10,5))
sns.heatmap(df.isnull(), cbar=False)
plt.title("Missing Data Heatmap")
plt.show()
```



Missing Value Imputation:

Two imputation techniques were applied. Statistical imputation using mean and mode was used for suitable columns. Additionally, KNN Imputation was implemented as an advanced method. After applying these techniques, the dataset was verified and all missing values were successfully handled.

IMPUTATION TECHNIQUES

Method 1 – Mean / Mode

```
df['Arrival Delay in Minutes'].fillna(df['Arrival Delay in Minutes'].mean(), inplace=True)
df['Gender'].fillna(df['Gender'].mode()[0], inplace=True)

/tmp/ipython-input-3222318373.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an in
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behave
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) inst

df['Arrival Delay in Minutes'].fillna(df['Arrival Delay in Minutes'].mean(), inplace=True)
/tmp/ipython-input-3222318373.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an in
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behave
For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) inst

df['Gender'].fillna(df['Gender'].mode()[0], inplace=True)
```

Method 2 – KNN Imputation

```
[7]
from sklearn.impute import KNNImputer

imputer = KNNImputer(n_neighbors=5)
df[['Departure Delay in Minutes']] = imputer.fit_transform(df[['Departure Delay in Minutes']])
```

Data Type Correction:

Necessary data type corrections were applied to ensure all numeric features were properly formatted. The date-related columns were parsed into datetime format for better usability.

DATA TYPE CORRECTION

```
ls
df['Flight Distance'] = pd.to_numeric(df['Flight Distance'], errors='coerce')
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 103904 entries, 0 to 103903
Data columns (total 25 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        103904 non-null   int64  
 1   id               103904 non-null   int64  
 2   Gender            103904 non-null   object  
 3   Customer Type    103904 non-null   object  
 4   Age               103904 non-null   int64  
 5   Type of Travel   103904 non-null   object  
 6   Class              103904 non-null   object  
 7   Flight Distance   103904 non-null   int64  
 8   Inflight wifi service 103904 non-null   int64  
 9   Departure/Arrival time convenient 103904 non-null   int64  
 10  Ease of Online booking 103904 non-null   int64  
 11  Gate location     103904 non-null   int64
```

Outlier Detection & Treatment:

Outliers in the Flight Distance column were detected using the Interquartile Range (IQR) method. Extreme values were removed in a documented and reproducible way to enhance dataset reliability.

OUTLIER DETECTION (IQR)&OUTLIER REMOVAL

```
[9] 0s
Q1 = df['Flight Distance'].quantile(0.25)
Q3 = df['Flight Distance'].quantile(0.75)

IQR = Q3 - Q1
lower = Q1 - 1.5 * IQR
upper = Q3 + 1.5 * IQR

df = df[(df['Flight Distance'] >= lower) & (df['Flight Distance'] <= upper)]
df.shape

(101613, 25)
```

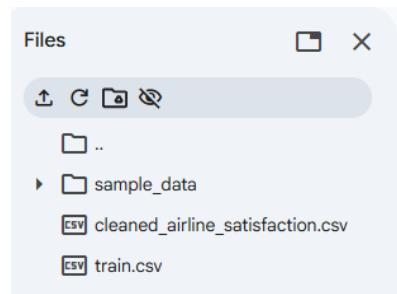
Final Clean Dataset:

FINAL DATA CHECK

```
[10] 0s
df.head()
df.tail()
```

| | Unnamed: 0 | id | Gender | Customer Type | Age | Type of Travel | Class | Flight Distance | Inflight wifi service | Departure/Arrival time convenient | ... | Inflight entertainment | On-board service | Leg room service | Baggage handling |
|--------|------------|-------|--------|-------------------|-----|-----------------|----------|-----------------|-----------------------|-----------------------------------|-----|------------------------|------------------|------------------|------------------|
| 103899 | 103899 | 94171 | Female | disloyal Customer | 23 | Business travel | Eco | 192 | 2 | 1 | ... | 2 | 3 | 1 | 4 |
| 103900 | 103900 | 73097 | Male | Loyal Customer | 49 | Business travel | Business | 2347 | 4 | 4 | ... | 5 | 5 | 5 | 5 |
| 103901 | 103901 | 68825 | Male | disloyal Customer | 30 | Business travel | Business | 1995 | 1 | 1 | ... | 4 | 3 | 2 | 4 |
| 103902 | 103902 | 54173 | Female | disloyal Customer | 22 | Business travel | Eco | 1000 | 1 | 1 | ... | 1 | 4 | 5 | 1 |

Saved CSV:



CLEANED CSV SAVE

```
[12]
df.to_csv("cleaned_airline_satisfaction.csv", index=False)
```